

# SIMPLE LINEAR REGRESSION

## ABOUT THE DATASET 'Salary Dataset'

### Variable Notes

1. Columns
2. YearsExperience
3. Salary

### DATA COLLECTION AND EXPLORATION

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [3]: # LOAD THE DATASET
data = pd.read_csv('Salary_dataset.csv')
data
```

Out[3]:

	Unnamed: 0	YearsExperience	Salary
0	0	1.2	39344.0
1	1	1.4	46206.0
2	2	1.6	37732.0
3	3	2.1	43526.0
4	4	2.3	39892.0
5	5	3.0	56643.0
6	6	3.1	60151.0
7	7	3.3	54446.0
8	8	3.3	64446.0
9	9	3.8	57190.0
10	10	4.0	63219.0
11	11	4.1	55795.0
12	12	4.1	56958.0
13	13	4.2	57082.0
14	14	4.6	61112.0
15	15	5.0	67939.0
16	16	5.2	66030.0
17	17	5.4	83089.0
18	18	6.0	81364.0
19	19	6.1	93941.0
20	20	6.9	91739.0
21	21	7.2	98274.0
22	22	8.0	101303.0
23	23	8.3	113813.0
24	24	8.8	109432.0
25	25	9.1	105583.0
26	26	9.6	116970.0
27	27	9.7	112636.0
28	28	10.4	122392.0
29	29	10.6	121873.0

In [4]:

```
# INSPECT THE DATA
data.head()
```

Out[4]:

	Unnamed: 0	YearsExperience	Salary
--	------------	-----------------	--------

0	0	1.2	39344.0
1	1	1.4	46206.0
2	2	1.6	37732.0
3	3	2.1	43526.0
4	4	2.3	39892.0

In [5]: `data.tail()`

Out[5]:

	Unnamed: 0	YearsExperience	Salary
--	------------	-----------------	--------

25	25	9.1	105583.0
26	26	9.6	116970.0
27	27	9.7	112636.0
28	28	10.4	122392.0
29	29	10.6	121873.0

In [6]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30 entries, 0 to 29
Data columns (total 3 columns):
#   Column             Non-Null Count  Dtype
---  -
0   Unnamed: 0         30 non-null     int64
1   YearsExperience    30 non-null     float64
2   Salary             30 non-null     float64
dtypes: float64(2), int64(1)
memory usage: 852.0 bytes
```

In [7]: `data.describe()`

Out[7]:


	Unnamed: 0	YearsExperience	Salary
--	------------	-----------------	--------

count	30.000000	30.000000	30.000000
mean	14.500000	5.413333	76004.000000
std	8.803408	2.837888	27414.429785
min	0.000000	1.200000	37732.000000
25%	7.250000	3.300000	56721.750000
50%	14.500000	4.800000	65238.000000
75%	21.750000	7.800000	100545.750000
max	29.000000	10.600000	122392.000000

In [58]: `plt.hist(data['YearsExperience'],bins=5)`


Out[58]:

```
(array([6., 9., 5., 4., 6.]),
 array([ 1.2,  3.08,  4.96,  6.84,  8.72, 10.6 ]),
 <BarContainer object of 5 artists>)
```




```
In [59]: data['Salary'].value_counts().plot(kind='pie', autopct='%.2f')
```

```
Out[59]: <Axes: ylabel='count'>
```



```
In [10]: plt.scatter(data['YearsExperience'], data['Salary'])
plt.xlabel("Years of Experience")
plt.ylabel("Salary")
```

```
Out[10]: Text(0, 0.5, 'Salary')
```



## CORRELATION

The code snippet removes the column 'Unnamed: 0' from the DataFrame 'data' and then computes the correlation matrix for the remaining columns using the `data.corr()` method.

```
In [13]: data.drop(['Unnamed: 0'], axis=1, inplace=True)
data.corr()
```

```
Out[13]:
```

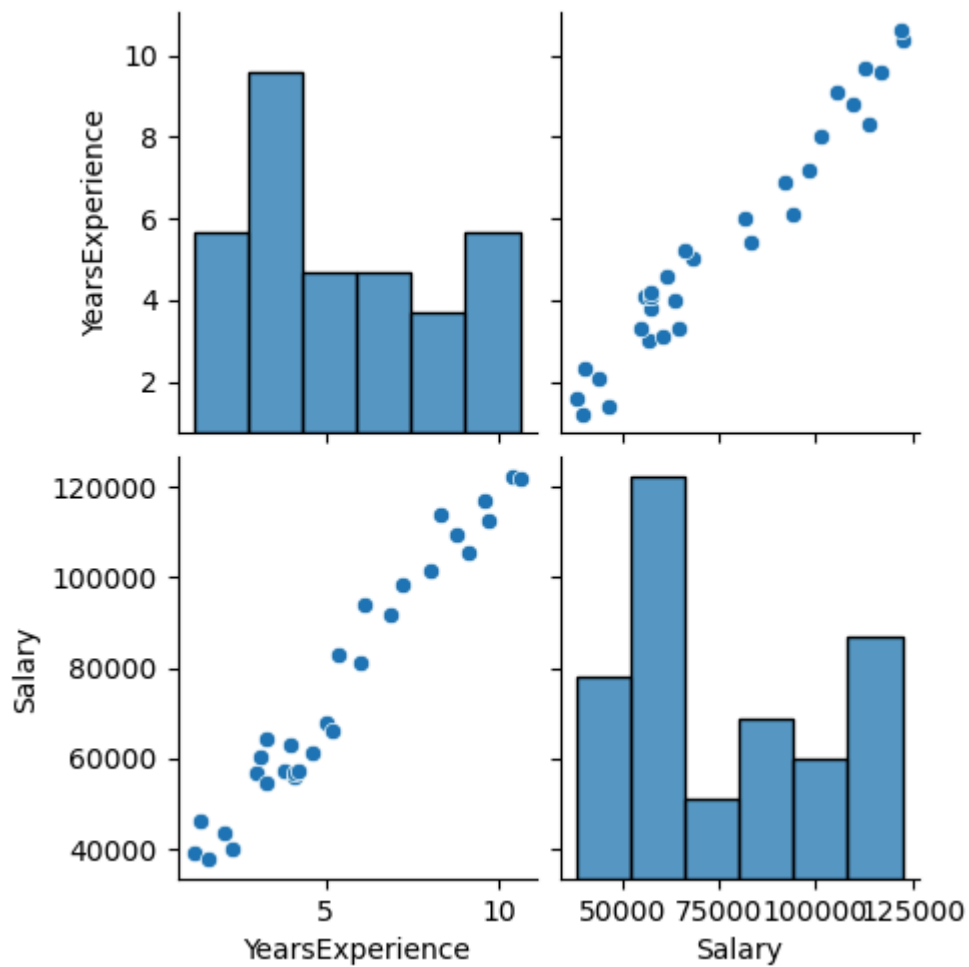
	YearsExperience	Salary
YearsExperience	1.000000	0.978242
Salary	0.978242	1.000000

## SEABORN FOR VISUALIZATION

```
In [16]: import seaborn as sns
sns.pairplot(data)
```

```
C:\Users\Nishita Bala\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
  self._figure.tight_layout(*args, **kwargs)
<seaborn.axisgrid.PairGrid at 0x2485b497f90>
```

```
Out[16]:
```



## INDEPENDENT AND DEPENDENT FEATURES

1. Independent feature should be a dataframe or 2-D array
2. Dependent variable can be in series form or 1-D array

```
In [19]: X = data[['YearsExperience']]
X
```

Out[19]:

YearsExperience	
0	1.2
1	1.4
2	1.6
3	2.1
4	2.3
5	3.0
6	3.1
7	3.3
8	3.3
9	3.8
10	4.0
11	4.1
12	4.1
13	4.2
14	4.6
15	5.0
16	5.2
17	5.4
18	6.0
19	6.1
20	6.9
21	7.2
22	8.0
23	8.3
24	8.8
25	9.1
26	9.6
27	9.7
28	10.4
29	10.6

In [22]: `np.array(X).shape`

Out[22]: (30, 1)

In [21]: `Y = data['Salary']`  
`Y`

```
Out[21]: 0      39344.0
         1      46206.0
         2      37732.0
         3      43526.0
         4      39892.0
         5      56643.0
         6      60151.0
         7      54446.0
         8      64446.0
         9      57190.0
        10      63219.0
        11      55795.0
        12      56958.0
        13      57082.0
        14      61112.0
        15      67939.0
        16      66030.0
        17      83089.0
        18      81364.0
        19      93941.0
        20      91739.0
        21      98274.0
        22     101303.0
        23     113813.0
        24     109432.0
        25     105583.0
        26     116970.0
        27     112636.0
        28     122392.0
        29     121873.0
        Name: Salary, dtype: float64
```

```
In [23]: np.array(Y).shape
```

```
Out[23]: (30,)
```

## TRAIN TEST SPLIT

The code uses `train_test_split` from `scikit-learn` to split dataset X and Y into training and testing sets. It assigns 25% to testing (`test_size=0.25`) and ensures reproducibility with `random_state=42`

```
In [24]: from sklearn.model_selection import train_test_split

         X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.25, random_state=42)
```

```
In [25]: X_train.shape
```

```
Out[25]: (22, 1)
```

```
In [26]: X_test.shape
```

```
Out[26]: (8, 1)
```

```
In [27]: Y_train.shape
```

```
Out[27]: (22,)
```

```
In [28]: Y_test.shape
```



Out[28]: (8,)

## STANDARDIZATION

The code initializes a StandardScaler to normalize data. It fits and transforms X\_train, scaling its features to zero mean and unit variance, and applies the same transformation to X\_test.

In [30]: `from sklearn.preprocessing import StandardScaler`

In [32]: `scaler = StandardScaler()  
X_train = scaler.fit_transform(X_train)`

In [34]: `X_test = scaler.transform(X_test)  
X_test`

C:\Users\Nishita Bala\anaconda3\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but StandardScaler was fitted with feature names

Out[34]: `warnings.warn(  
array([[-1.15872417],  
[-1.81577634],  
[-1.35444184],  
[-1.759857 ],  
[-2.0534335 ],  
[-1.98353433],  
[-1.06086534],  
[-1.28454267]])`

## APPLY LINEAR REGRESSION

The code imports the LinearRegression model from sklearn.linear\_model, initializes it with n\_jobs=-1 for parallel processing, and fits the model to training data X\_train and Y\_train.

In [35]: `from sklearn.linear_model import LinearRegression`

In [38]: `regression = LinearRegression(n_jobs = -1)`

In [39]: `regression.fit(X_train,Y_train)`

Out[39]: `LinearRegression  
LinearRegression(n_jobs=-1)`


In [42]: `print("Coefficient or slope (Beta1):",regression.coef_)  
print("Intercept (Beta0):",regression.intercept_)`

Coefficient or slope (Beta1): [25063.1519945]  
Intercept (Beta0): 70417.40909090909

## PLOT TRAINING DATA PLOT BEST FIT LINE

In [44]: `plt.scatter(X_train,Y_train)  
plt.plot(X_train,regression.predict(X_train))`

Out[44]: `[<matplotlib.lines.Line2D at 0x2485d2d3590>]`



## PREDICTION FOR TEST DATA

```
In [71]: Y_pred = regression.predict(X_test)
         y_pred
```

```
Out[71]: array([41376.1290096 , 24908.33081684, 36470.82742027, 26309.84555665,
                18951.89317266, 20703.78659742, 43828.77980427, 38222.72084503])
```

## PERFORMANCE MATRICS

```
In [69]: from sklearn.metrics import mean_absolute_error, mean_squared_error
```

```
In [70]: mse = mean_squared_error(Y_test,Y_pred)
         mae = mean_absolute_error(Y_test,Y_pred)
         rmse = np.sqrt(mse)

         print("Mean Square Error",mse)
         print("Mean Absolute",mae)
         print("Root Mean Square Error",rmse)
```

Mean Square Error 3847398784.0427675

Mean Absolute 60020.58584715701

Root Mean Square Error 62027.40349267223